

## Section 1

Microarchitecture Side-Channel Resistant  
Instructions Spans (scripsans)

## Examples semantics

- ▶ `spansec.create [policies]`: create a new scrispan with given optional set of security policies (flags).
- ▶ `spansec.save rd`: save the current scrispan configuration (ID + security policy) in a register.
- ▶ `spansec.restore rs1`: restore a scrispan configuration from a register holding ID and security policy.

**We assume that a change of scrispan ID implies microarchitectural isolation.**

- ▶ `spansec.alter [policies]`: in addition we could add an instruction that keeps the ID unchanged but modify the security policies.

**We are not discussing the implementation for now. Have you alternative semantics in mind ?**

# Security policies

[policies] is the set of security policies applied to a scripsan.

- ▶ no speculation
- ▶ inline memory encryption
- ▶ constant time execution

*Personal comment (Ronan): security policies add complexity, but not sure if useful. Suggested policies could be better achieved with other means.*

**Do you have strong use cases for security policies ?**

## Usecase #1: web server

Goal: isolate users

Each process has its own scrispan.

*// creating process P1*

`spansec.create`

*// during switch from P1 to P2*

`spansec.save P1`

`spansec.restore P2`

### Comments

- ▶ ✓ This explicit span gymnastic ensures proper microarchitectural state isolation.
- ▶ ✗ In this precise use case, we could tie the scrispan to the ASID (Address Space ID) instead.

## Pattern: Microarchitectural state poisoning

Goal: isolating user poisoned uarch state from data processing

*// span A: user controlled, span B: data processing*

```
spansec.restore A  
state_poisoning();  
spansec.save A
```

```
spansec.restore B  
data_processing();  
spansec.save B
```

### Comments

Since poisoning can be done in the same or another address space (cf <https://transient.fail>), we cannot rely on tying the scripsan with the ASID.

# Simultaneous multithreading (SMT) I

Example from [https://github.com/IAIK/transientfail/blob/master/pocs/spectre/RSB/sa\\_ip/main.c](https://github.com/IAIK/transientfail/blob/master/pocs/spectre/RSB/sa_ip/main.c), two software threads in the same address space.

## Pattern

```
pthread_t attacker_thread;  
pthread_t victim_thread;  
pthread_create(&attacker_thread, 0, attacker, 0);  
pthread_create(&victim_thread, 0, victim, 0);
```

## Solution

```
// pthread_create should start with  
spansec.create
```

# Simultaneous multithreading (SMT) II

## Hardware consequences

- ▶ They are scheduled on different cores.
- ▶ If the hardware allows, the harts are isolated in the same core.
- ▶ They are not executed in parallel but sequentially, with proper microarchitectural flushing when switching between threads.

# Spectre-PHT (v1) I

## Vulnerability

```
if (x < array1_size) {  
    y = array2[array1[x] * 4096];  
}  
  
//speculative load 1: array1[x] -> accessing secret  
//speculative load 2: array2[...] -> leaking gadget
```

## Solution 1: disabling speculation

```
spansec.alter +nospeculation  
if (x < array1_size) {  
    y = array2[array1[x] * 4096];  
}  
spansec.alter -nospeculation
```



# Spectre-PHT (v1) II

## Solution 2: isolating poisoned state

```
if (x < array1_size) {  
    spansec.save A  
    spansec.create  
    y = array2[array1[x] * 4096];  
    spansec.restore A  
}
```

## Comments

- ▶ **X** nospeculation: poor portability, precise behaviour is very hardware dependent, overkill ? (How to allow speculation across scripan ?).
- ▶ **X** state isolation: very bad performances.
- ▶ Probably needs another mechanism.